



# MapReduce 2.0程序设计 (涉及多语言)-理论部分

讲师：董西成

博客：[dongxicheng.org](http://dongxicheng.org)

微信号：[hadoop-123](#)

(二维码见右)



Open Passion Value



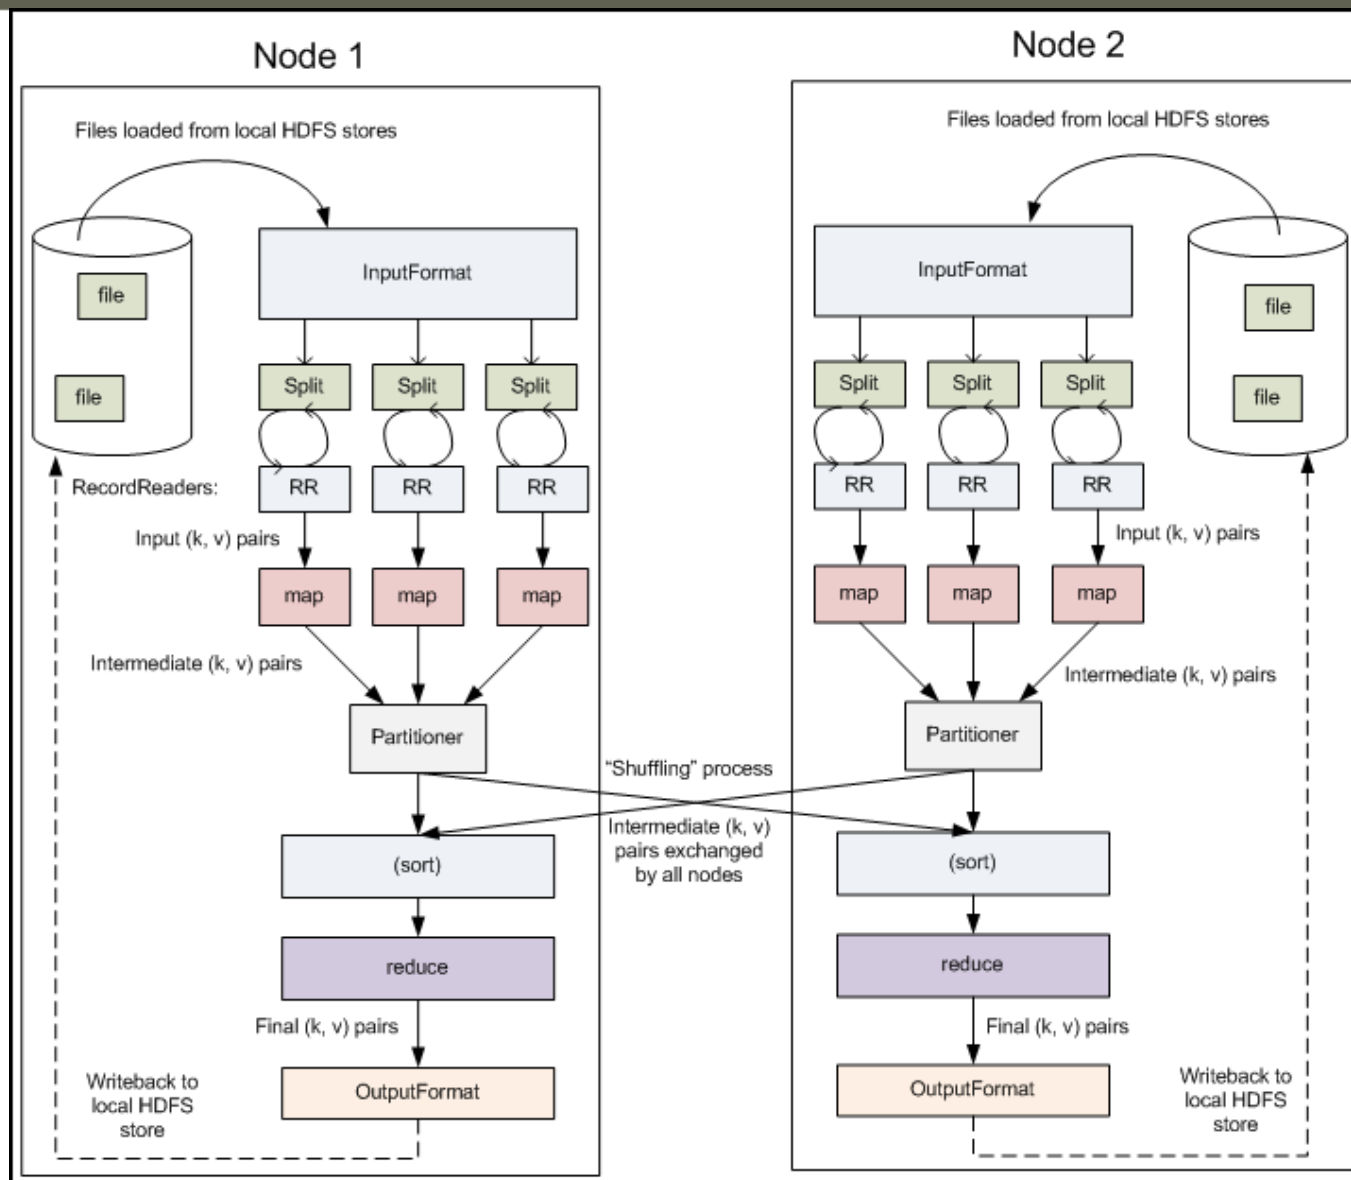


1. MapReduce 2.0的编程模型
2. MapReduce 2.0编程接口介绍
3. Java编程
4. 多语言编程
5. 总结



- **MapReduce**将整个运行过程分为两个阶段：**Map**阶段和**Reduce**阶段
- **Map**阶段由一定数量的**Map Task**组成
  - ✓ 输入数据格式解析：**InputFormat**
  - ✓ 输入数据处理：**Mapper**
  - ✓ 数据分组：**Partitioner**
- **Reduce**阶段由一定数量的**Reduce Task**组成
  - ✓ 数据远程拷贝
  - ✓ 数据按照**key**排序
  - ✓ 数据处理：**Reducer**
  - ✓ 数据输出格式：**OutputFormat**

# MapReduce编程模型—外部物理结构





## ➤ Map阶段

- ✓ InputFormat （默认TextInputFormat）

- ✓ Mapper

- ✓ Combiner （local reducer）

- ✓ Partitioner

## ➤ Reduce阶段

- ✓ Reducer

- ✓ OutputFormat （默认TextOutputFormat）



1. MapReduce的编程模型
2. MapReduce编程接口介绍
3. Java编程
4. 多语言编程
5. 总结



- **Hadoop**提供了三种编程方式：
  - ✓ **Java**（最原始的方式）
  - ✓ **Hadoop Streaming**（支持多语言）
  - ✓ **Hadoop Pipes**（支持C/C++）
- **Java**编程接口是所有编程方式的基础；
- 不同的编程接口只是暴露给用户的形式不同而已，内部执行引擎是一样的；
- 不同编程方式效率不同。





- **Java编程接口组成；**
  - ✓ 旧API：所在java包：`org.apache.hadoop.mapred`
  - ✓ 新API：所在java包：`org.apache.hadoop.mapreduce`
- **新API具有更好的扩展性；**
- **两种编程接口只是暴露给用户的形式不同而已，内部执行引擎是一样的；**
- **旧API可以完全兼容Hadoop 2.0，但新API不行。**





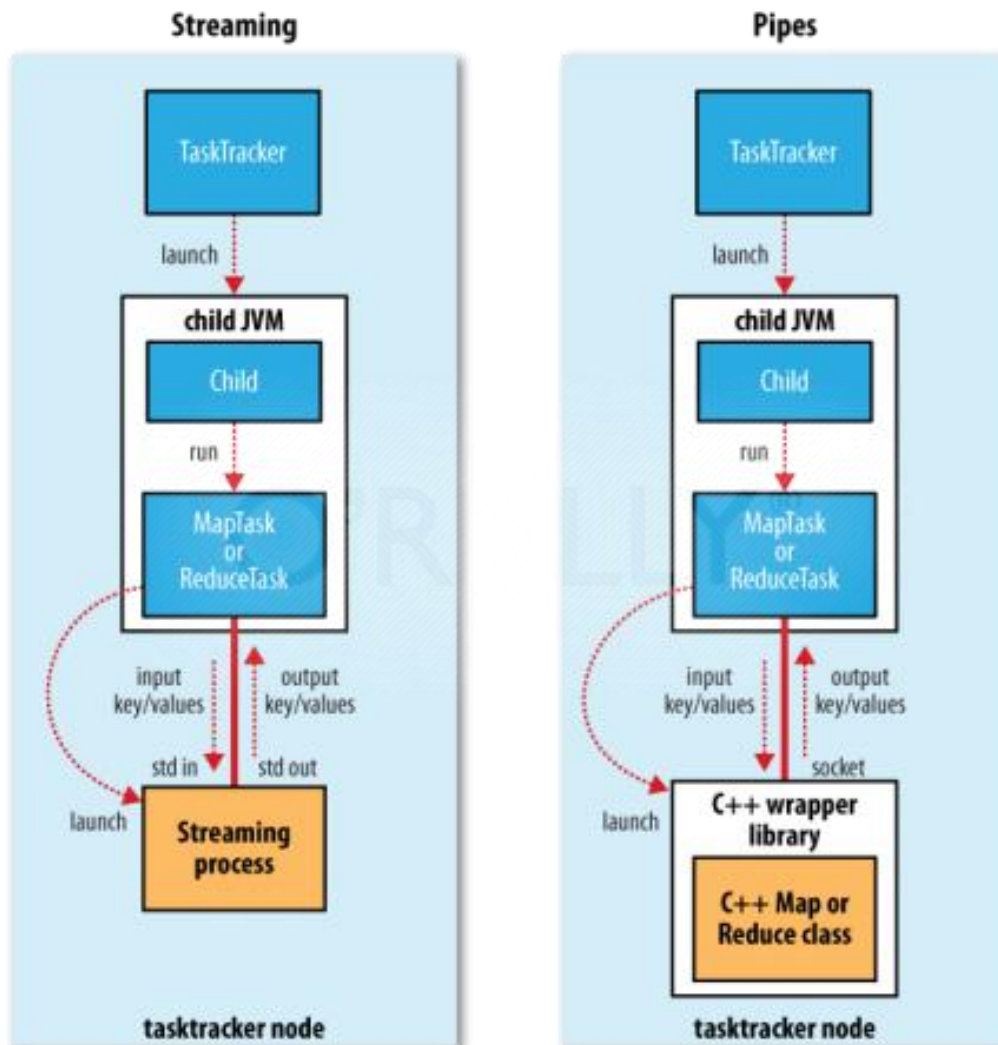
## ➤ 从hadoop 1.0.0开始，所有发行版均包含新旧两类API;

- hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-core/src/main/java
  - org.apache.hadoop.filecache
  - org.apache.hadoop.mapred
  - org.apache.hadoop.mapred.jobcontrol
  - org.apache.hadoop.mapred.join
  - org.apache.hadoop.mapred.lib
  - org.apache.hadoop.mapred.lib.aggregate
  - org.apache.hadoop.mapred.lib.db
  - org.apache.hadoop.mapred.pipes
  - org.apache.hadoop.mapreduce
  - org.apache.hadoop.mapreduce.counters
  - org.apache.hadoop.mapreduce.filecache
  - org.apache.hadoop.mapreduce.jobhistory
  - org.apache.hadoop.mapreduce.lib.aggregate
  - org.apache.hadoop.mapreduce.lib.chain
  - org.apache.hadoop.mapreduce.lib.db



- 与Linux管道机制一致
- 通过标准输入输出实现进程间通信
- 标准输入输出是任何语言都有的
- 几个举例：
  - ✓ `cat 1.txt | grep “dong” | sort`
  - ✓ `cat 1.txt | python grep.py | java sort.jar`

# Hadoop Streaming/pipes





1. MapReduce的编程模型
2. MapReduce编程接口介绍
3. Java编程
4. 多语言编程
5. 总结

# 实例1: WordCount问题



## Input:

File containing words

```
Hello World Bye World  
Hello Hadoop Bye Hadoop  
Bye Hadoop Hello Hadoop
```

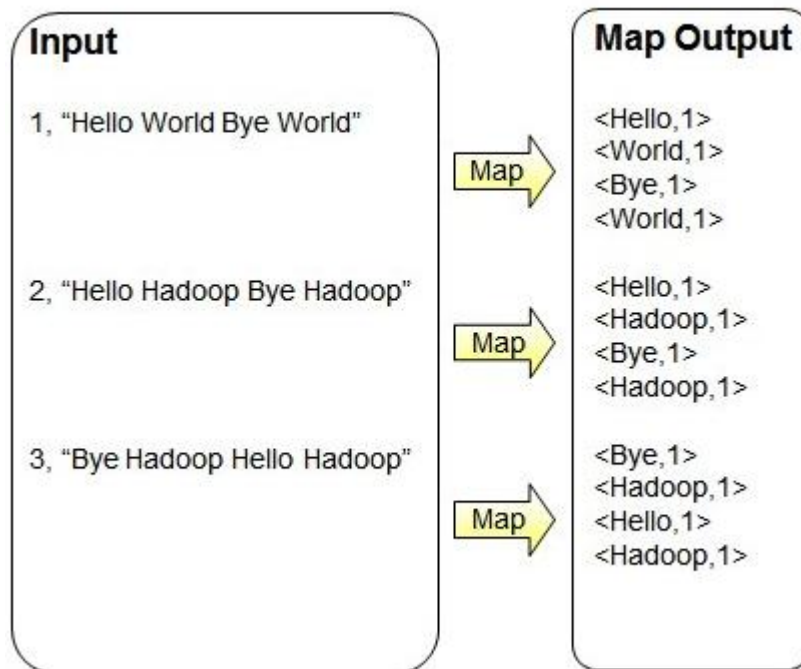
MapReduce

## Output:

Number of occurrences  
of each word

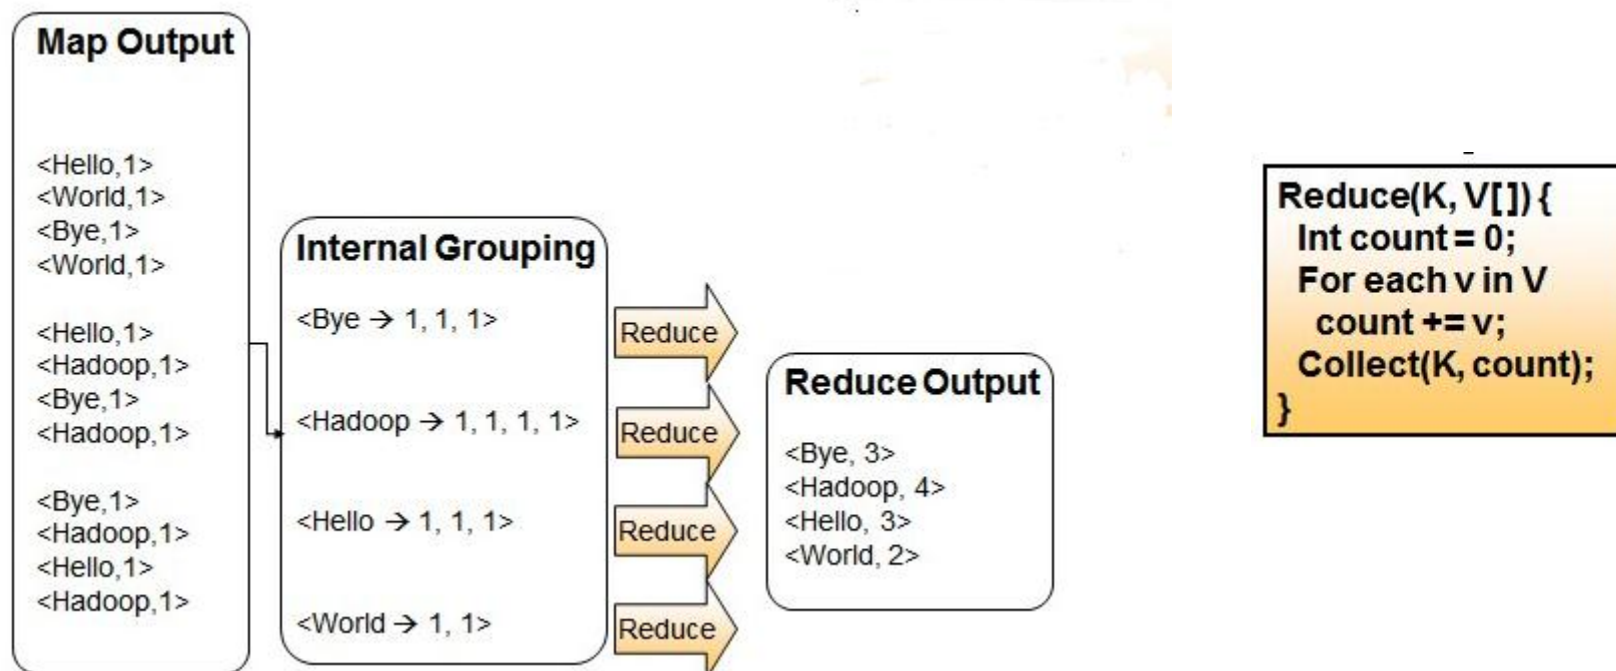
```
Bye 3  
Hadoop 4  
Hello 3  
World 2
```

# WordCount问题—map阶段



```
Map(K, V) {  
  For each word w in V  
    Collect(w, 1);  
}
```

# WordCount问题—reduce阶段





# WordCount问题—mapper设计与实现



```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

# WordCount问题—reducer设计与实现



```
public static class IntSumReducer extends
    Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

# WordCount问题—main函数设计与实现



```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(conf, args)
        .getRemainingArgs();
    if (otherArgs.length != 2) {
        System.err.println("Usage: wordcount <in> <out>");
        System.exit(2);
    }
    Job job = new Job(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
    FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

# WordCount问题—程序运行



```
yarn@SY-0266:/opt/pgs/yarn-client$ bin/hadoop fs -mkdir /test/input
yarn@SY-0266:/opt/pgs/yarn-client$ bin/hadoop fs -put streaming-examples/hadoop-hdfs-namenode-SY-0245.log /test/input
yarn@SY-0266:/opt/pgs/yarn-client$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-2.2.0.jar wordcount /test/input /test/output
14/03/01 13:50:43 INFO client.RMPProxy: Connecting to ResourceManager at /10.10.65.13:8032
14/03/01 13:50:43 INFO input.FileInputFormat: Total input paths to process : 1
14/03/01 13:50:43 INFO mapreduce.JobSubmitter: number of splits:1
14/03/01 13:50:43 INFO Configuration.deprecation: user.name is deprecated. Instead, use mapred
```

```
14/03/01 13:50:45 INFO mapreduce.Job: Running job: job_1393577861371_0005
14/03/01 13:50:49 INFO mapreduce.Job: Job job_1393577861371_0005 running in uber mode : false
14/03/01 13:50:49 INFO mapreduce.Job: map 0% reduce 0%
14/03/01 13:50:55 INFO mapreduce.Job: map 100% reduce 0%
14/03/01 13:51:01 INFO mapreduce.Job: map 100% reduce 100%
14/03/01 13:51:01 INFO mapreduce.Job: Job job_1393577861371_0005 completed successfully
14/03/01 13:51:01 INFO mapreduce.Job: Counters: 43
    File System Counters
      FILE: Number of bytes read=69812
      FILE: Number of bytes written=307709
      FILE: Number of read operations=0
      FILE: Number of large read operations=0
      FILE: Number of write operations=0
      HDFS: Number of bytes read=750854
      HDFS: Number of bytes written=60778
```

# WordCount问题—程序运行

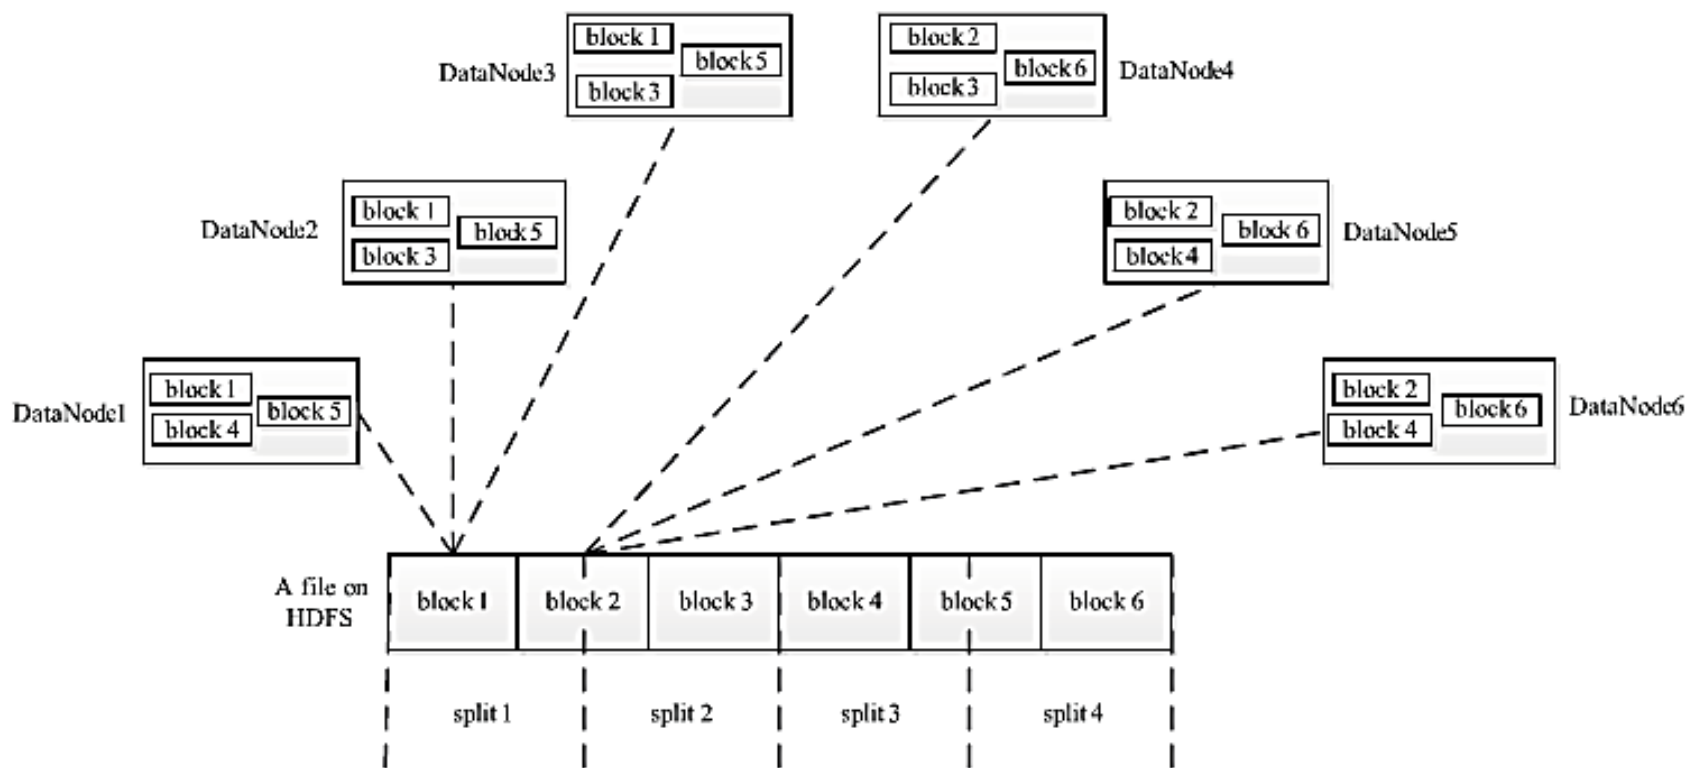


```
yarn@SY-0266:/opt/pgs/yarn-client$ bin/hadoop fs -cat /test/output/part-r-00000 | head
#          16
#1         12
#11        1
#12        1
#14        1
#16        1
#19        1
#3         2
#5         2
#7         2
```



- 使用默认的TextInputFormat
  - ✓ 每个Map Task处理一个split;
  - ✓ 一个split大小等于一个block;
  - ✓ 如果最后一行数据被截断，则读取后一个block前半部分;
  - ✓ 转换成key/value对，key是偏移量，value是行内容。

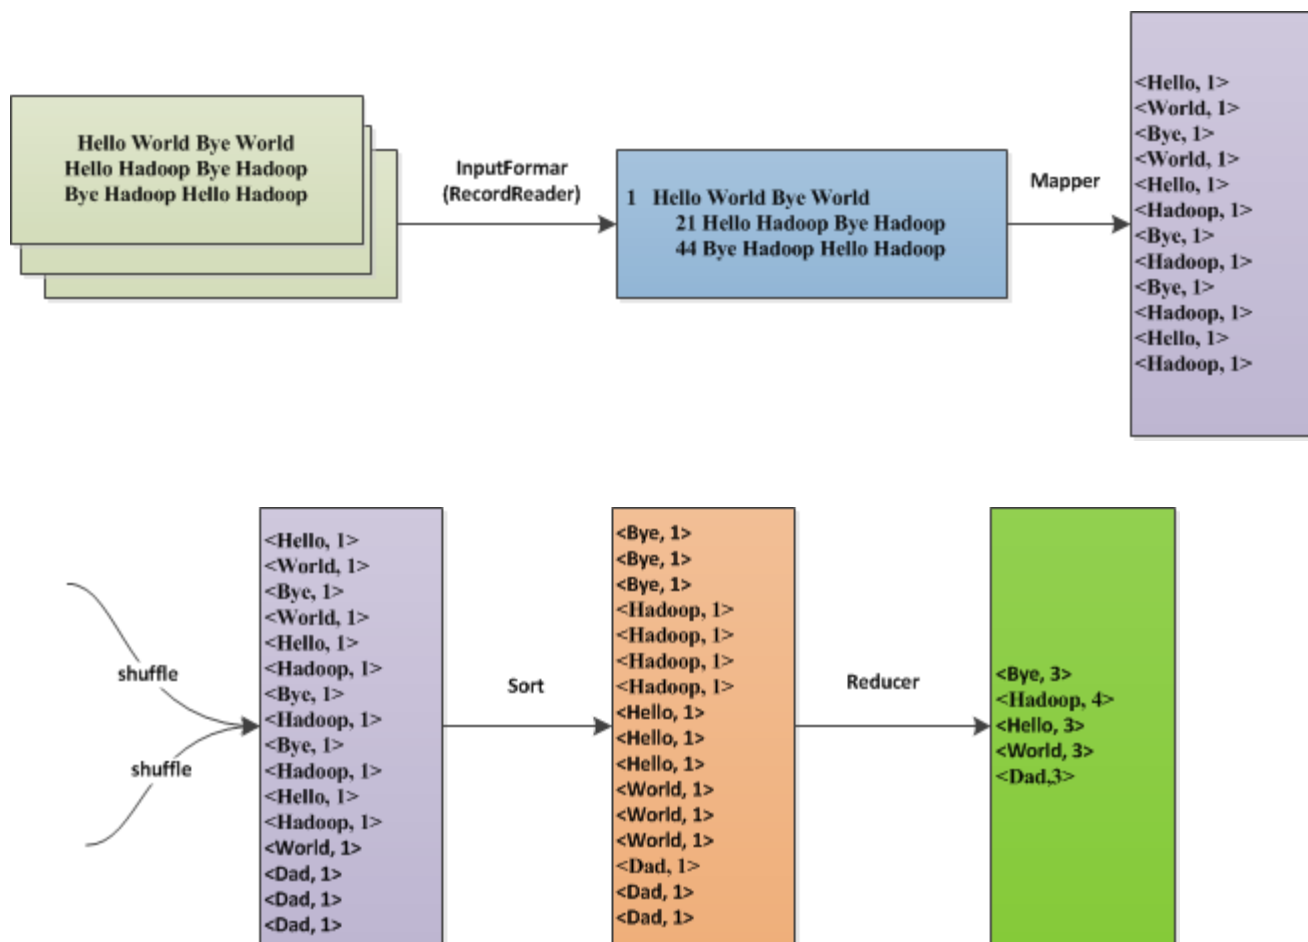
# Wordcount问题—输入数据格式解析



split 与 block 的对应关系



# Wordcount问题—数据流



# InputFormat—输入数据解析



```
public interface InputFormat<K, V> {
```

```
    InputSplit[] getSplits(JobConf job, int numSplits) throws IOException;
```

```
    RecordReader<K, V> getRecordReader(InputSplit split,  
                                         JobConf job,  
                                         Reporter reporter) throws IOException;
```

```
}
```

- 默认为TextInputFormat，针对文本文件的；
- 用户可通过参数mapred.input.format.class设置InputFormat实现

# Mapper—map处理逻辑



**org.apache.hadoop.mapred (旧API) :**

```
public interface Mapper<K1, V1, K2, V2> extends JobConfigurable, Closeable {  
  
    void map(K1 key, V1 value, OutputCollector<K2, V2> output, Reporter reporter)  
    throws IOException;  
  
}
```

- **新API位于org.apache.hadoop.mapreduce.Mapper中;**
- **新API更加灵活。**

# Partitioner—map输出结构分片



**org.apache.hadoop.mapred (旧API) :**

```
public interface Partitioner<K2, V2> extends JobConfigurable {  
  
    int getPartition(K2 key, V2 value, int numPartitions);  
}
```

**org.apache.hadoop.mapreduce (新API) :**

```
public abstract class Partitioner<KEY, VALUE> {  
  
    public abstract int getPartition(KEY key, VALUE value, int numPartitions);  
}  
  
public class HashPartitioner<K2, V2> implements Partitioner<K2, V2> {  
  
    public void configure(JobConf job) {}  
  
    /** Use {@link Object#hashCode()} to partition. */  
    public int getPartition(K2 key, V2 value,  
                        int numReduceTasks) {  
        return (key.hashCode() & Integer.MAX_VALUE) % numReduceTasks;  
    }  
}
```

# Reducer—reduce处理逻辑

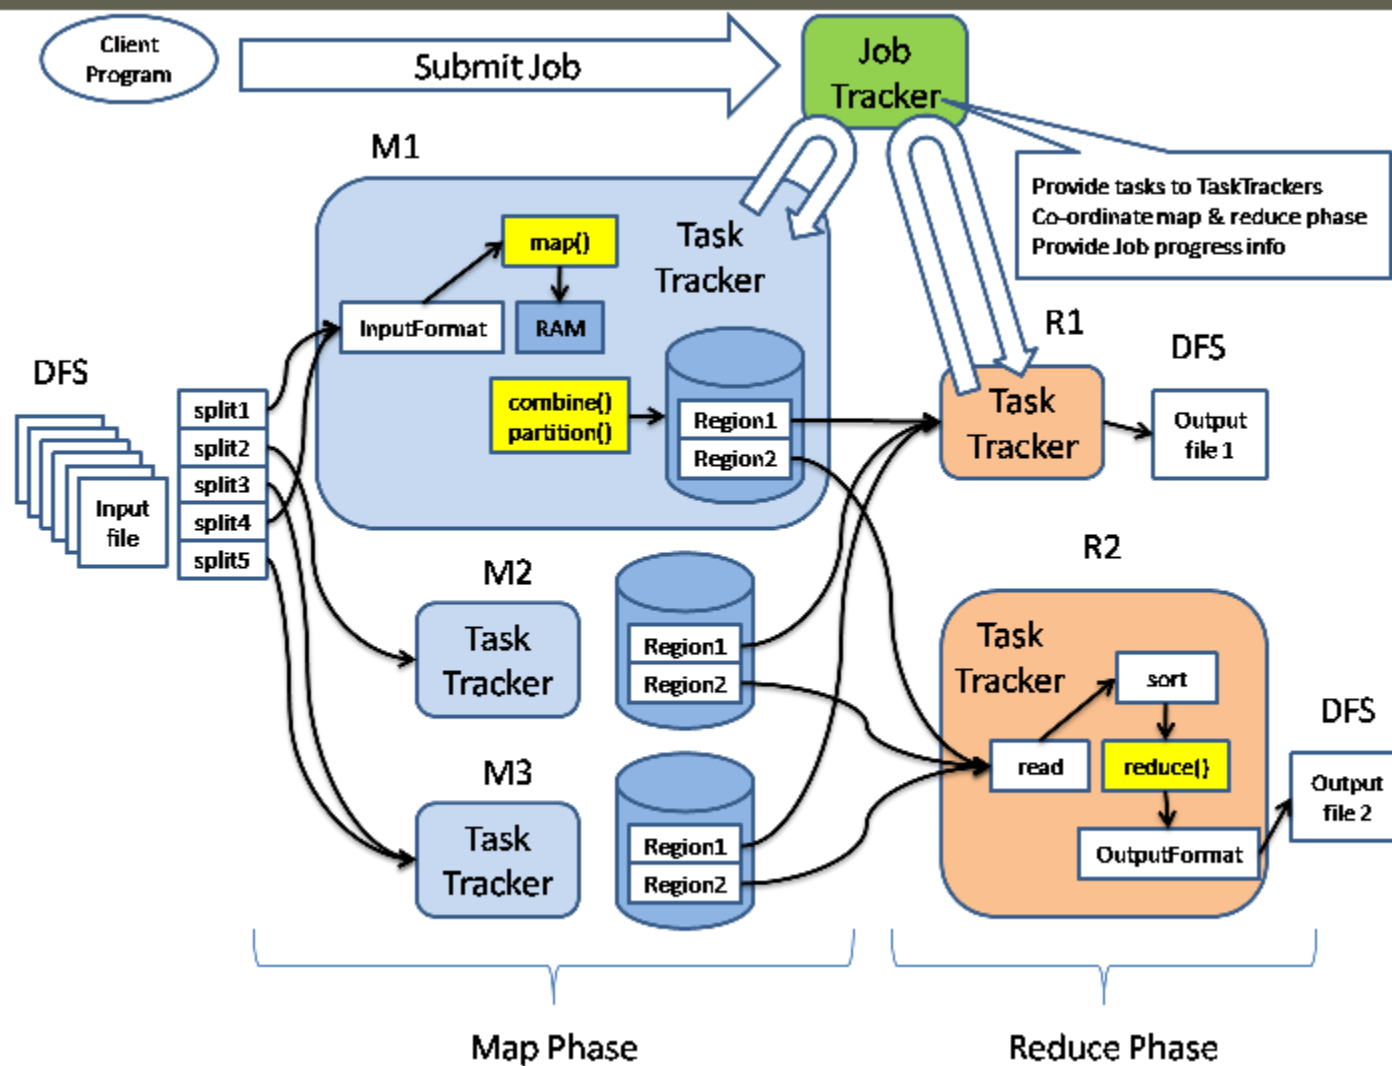


**org.apache.hadoop.mapred (旧API) :**

```
public interface Reducer<K2, V2, K3, V3> extends JobConfigurable, Closeable {  
  
    void reduce(K2 key, Iterator<V2> values,  
               OutputCollector<K3, V3> output, Reporter reporter)  
        throws IOException;  
}
```

- **新API位于org.apache.hadoop.mapreduce.Reducer中;**
- **新API更加灵活。**

# 小结



## 实例2: Grep问题




一批TB或者PB量级的文档，需要完成以下功能：

- 搜索符合某种规则（正则表达式）的单词或者句子；
- 统计相应的单词或者句子的数目；
- 按照数目对其进行排序，并输出最终结果。




# Grep问题解决思路



HDFS中的文档


shop[A-Z.]+



shopA	10
shopC	25
shopB	5
shopD	100

频率统计结果

排序



shopB	5
shopA	10
shopC	25
shopD	100

排序结果



分为两个作业：

- 作业1： WordCount

- ✓统计符合某种规则的单词数目

- 作业2： Sort

- ✓按照单词数目进行全排序

- ✓依赖于前一个作业的输出结果

# Grep问题解决思路



```
/* Extracts matching regexs from input files and counts them. */
public class Grep extends Configured implements Tool {
    private Grep() {} // singleton

    public int run(String[] args) throws Exception {
        if (args.length < 3) {
            System.out.println("Grep <inDir> <outDir> <regex> [<group>]");
            ToolRunner.printGenericCommandUsage(System.out);
            return -1;
        }

        Path tempDir =
            new Path("grep-temp-"+
                Integer.toString(new Random().nextInt(Integer.MAX_VALUE)));

        JobConf grepJob = new JobConf(getConf(), Grep.class);

        try {

            grepJob.setJobName("grep-search");

            FileInputFormat.setInputPaths(grepJob, args[0]);

            grepJob.setMapperClass(RegexMapper.class);
            grepJob.set("mapred.mapper.regex", args[2]);
            if (args.length == 4)
                grepJob.set("mapred.mapper.regex.group", args[3]);

            grepJob.setCombinerClass(LongSumReducer.class);
            grepJob.setReducerClass(LongSumReducer.class);

            FileOutputFormat.setOutputPath(grepJob, tempDir);
            grepJob.setOutputFormat(SequenceFileOutputFormat.class);
            grepJob.setOutputKeyClass(Text.class);
            grepJob.setOutputValueClass(LongWritable.class);

            JobClient.runJob(grepJob);
        }
    }
}
```

# Grep问题解决思路



```
JobConf sortJob = new JobConf(Grep.class);
sortJob.setJobName("grep-sort");

FileInputFormat.setInputPaths(sortJob, tempDir);
sortJob.setInputFormat(SequenceFileInputFormat.class);

sortJob.setMapperClass(InverseMapper.class);

sortJob.setNumReduceTasks(1);           // write a single file
FileOutputFormat.setOutputPath(sortJob, new Path(args[1]));
sortJob.setOutputKeyComparatorClass      // sort by decreasing freq
(LongWritable.DecreasingComparator.class);

JobClient.runJob(sortJob);
}
finally {
    FileSystem.get(grepJob).delete(tempDir, true);
}
return 0;
}

public static void main(String[] args) throws Exception {
    int res = ToolRunner.run(new Configuration(), new Grep(), args);
    System.exit(res);
}
```

# Grep问题解决思路



```
/** A {@link Mapper} that extracts text matching a regular expression. */
public class RegexMapper<K> extends MapReduceBase
    implements Mapper<K, Text, Text, LongWritable> {

    private Pattern pattern;
    private int group;

    public void configure(JobConf job) {
        pattern = Pattern.compile(job.get("mapred.mapper.regex"));
        group = job.getInt("mapred.mapper.regex.group", 0);
    }

    public void map(K key, Text value,
        OutputCollector<Text, LongWritable> output,
        Reporter reporter)
        throws IOException {
        String text = value.toString();
        Matcher matcher = pattern.matcher(text);
        while (matcher.find()) {
            output.collect(new Text(matcher.group(group)), new LongWritable(1));
        }
    }
}
```

# Grep问题解决思路



```
public class InverseMapper<K, V>
    extends MapReduceBase implements Mapper<K, V, V, K> {

    /** The inverse function.  Input keys and values are swapped.*/
    public void map(K key, V value,
        OutputCollector<V, K> output, Reporter reporter)
        throws IOException {
        output.collect(value, key);
    }
}
```

# Grep程序运行



```
yarn@SY-0266:/opt/pgs/yarn-client$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-2.2.0.jar grep /test/input /test/output2 software*
14/03/01 13:54:26 INFO client.RMProxy: Connecting to ResourceManager at /10.10.65.13:8032
14/03/01 13:54:26 WARN mapreduce.JobSubmitter: No job jar file set. User classes may not be found. See Job or Job#setJar(String).
14/03/01 13:54:26 INFO input.FileInputFormat: Total input paths to process : 1
14/03/01 13:54:26 INFO mapreduce.JobSubmitter: number of splits:1
14/03/01 13:54:26 INFO Configuration.deprecation: user.name is deprecated. Instead, use mapreduce.job.user.name
14/03/01 13:54:26 INFO Configuration.deprecation: mapred.output.value.class is deprecated. Instead, use mapreduce.job.output.value.class
14/03/01 13:54:26 INFO Configuration.deprecation: mapreduce.combine.class is deprecated. Instead, use mapreduce.job.combine.class
14/03/01 13:54:26 INFO Configuration.deprecation: mapreduce.map.class is deprecated. Instead, use mapreduce.job.map.class
14/03/01 13:54:26 INFO Configuration.deprecation: mapred.job.name is deprecated. Instead, use mapreduce.job.name
14/03/01 13:54:26 INFO Configuration.deprecation: mapreduce.reduce.class is deprecated. Instead, use mapreduce.job.reduce.class
14/03/01 13:54:26 INFO Configuration.deprecation: mapred.input.dir is deprecated. Instead, use mapreduce.input.fileinputformat.inputdir
14/03/01 13:54:26 INFO Configuration.deprecation: mapred.output.dir is deprecated. Instead, use mapreduce.output.fileoutputformat.outputdir
14/03/01 13:54:26 INFO Configuration.deprecation: mapreduce.outputformat.class is deprecated. Instead, use mapreduce.job.outputformat.class
14/03/01 13:54:26 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
14/03/01 13:54:26 INFO Configuration.deprecation: mapred.output.key.class is deprecated. Instead, use mapreduce.job.output.key.class
14/03/01 13:54:26 INFO Configuration.deprecation: mapred.working.dir is deprecated. Instead, use mapreduce.job.working.dir
14/03/01 13:54:26 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1393577861371_0006
14/03/01 13:54:26 INFO mapred.YARNRunner: Job jar is not present. Not adding any jar to the list of resources.
14/03/01 13:54:28 INFO impl.YarnClientImpl: Submitted application application_1393577861371_0006 to ResourceManager at /10.10.65.13:8032
14/03/01 13:54:28 INFO mapreduce.Job: The url to track the job: http://SY-0244:8088/proxy/application_1393577861371_0006/
14/03/01 13:54:28 INFO mapreduce.Job: Running job: job_1393577861371_0006
14/03/01 13:54:33 INFO mapreduce.Job: Job job_1393577861371_0006 running in uber mode : false
14/03/01 13:54:33 INFO mapreduce.Job: map 0% reduce 0%
14/03/01 13:54:38 INFO mapreduce.Job: map 100% reduce 0%
14/03/01 13:54:43 INFO mapreduce.Job: map 100% reduce 100%
14/03/01 13:54:43 INFO mapreduce.Job: Job job_1393577861371_0006 completed successfully
14/03/01 13:54:43 INFO mapreduce.Job: Counters: 43
  File System Counters
    FILE: Number of bytes read=6
    FILE: Number of bytes written=168471
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
```



# Grep程序运行



```
14/03/01 13:54:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1393577861371_0007
14/03/01 13:54:43 INFO mapred.YARNRunner: Job jar is not present. Not adding any jar to the list of resources.
14/03/01 13:54:44 INFO impl.YarnClientImpl: Submitted application application_1393577861371_0007 to ResourceManager at /10.10.65.13:8032
14/03/01 13:54:44 INFO mapreduce.Job: The url to track the job: http://SY-0244:8088/proxy/application_1393577861371_0007/
14/03/01 13:54:44 INFO mapreduce.Job: Running job: job_1393577861371_0007
14/03/01 13:54:48 INFO mapreduce.Job: Job job_1393577861371_0007 running in uber mode : false
14/03/01 13:54:48 INFO mapreduce.Job: map 0% reduce 0%
14/03/01 13:54:53 INFO mapreduce.Job: map 100% reduce 0%
14/03/01 13:54:58 INFO mapreduce.Job: map 100% reduce 100%
14/03/01 13:54:59 INFO mapreduce.Job: Job job_1393577861371_0007 completed successfully
14/03/01 13:54:59 INFO mapreduce.Job: Counters: 43
  File System Counters
    FILE: Number of bytes read=6
    FILE: Number of bytes written=167305
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=213
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=7
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=5254
    Total time spent by all reduces in occupied slots (ms)=5794
  Map-Reduce Framework
    Map input records=0
    Map output records=0
    Map output bytes=0
    Map output materialized bytes=6
    Input split bytes=127
    Combine input records=0
    Combine output records=0
    Reduce input groups=0
    Reduce shuffle bytes=6
```



1. MapReduce的编程模型
2. MapReduce编程接口介绍
3. Java编程
4. 多语言编程
5. 总结



- 以标准输入流作为输入；
  - ✓ C++: `cin`
  - ✓ C: `scanf`
- 以标准输出流作为输出；
  - ✓ C++: `cout`
  - ✓ C: `printf`
- 可实现Mapper和Reducer，其他组件（  
InputFormat、Partitioner等需要用Java语言实现）

# 实例1：用C++实现Wordcount（Mapper实现）



```
#include <iostream>
#include <string>
using namespace std;

int main() {
    string key;
    while(cin >> key) {
        cout << key << "\t" << "1" << endl;
    }
    return 0;
}
```

# 实例1：用C++实现Wordcount（Reducer实现）



```
#include <iostream>
#include <string>

using namespace std;
int main() {
    string cur_key, last_key, value;
    cin >> cur_key >> value;
    last_key = cur_key;
    int n = 1;
    while(cin >> cur_key) {
        cin >> value;
        if(last_key != cur_key) {
            cout << last_key << "\t" << n << endl;
            last_key = cur_key;
            n = 1;
        } else {
            n++;
        }
    }
    cout << last_key << "\t" << n << endl;
    return 0;
}
```

# 实例1：用C++实现Wordcount（编译运行）



- 编译程序，生成可执行文件；

- ✓ `g++ -o mapper mapper.cpp`

- ✓ `g++ -o reducer reducer.cpp`

- 测试程序；

- ✓ `cat test.txt | ./mapper | sort`

```
root@SY-0266:/opt/pgs/yarn-client/streaming-examples# cat test.txt
```

```
i  
have  
a  
book  
you  
do  
not  
have  
one  
so  
i  
am  
better  
than  
you  
ha  
ha
```

```
root@SY-0266:/opt/pgs/yarn-client/streaming-examples# cat test.txt | ./mapper | sort | ./reducer
```

```
a      1  
am     1  
better 1  
book   1  
do     1  
ha     2  
have   2  
i      2  
not    1  
one    1  
so     1  
than   1  
you    2
```

# 实例1：用C++实现Wordcount（编译运行）



## ● Hadoop上运行Wordcount程序

```
#!/bin/bash
HADOOP_HOME=/opt/pgs/yarn-client
INPUT_PATH=/test/input
OUTPUT_PATH=/test/output
echo "Clearing output path: $OUTPUT_PATH"
$HADOOP_HOME/bin/hadoop fs -rmr $OUTPUT_PATH

${HADOOP_HOME}/bin/hadoop jar\
    ${HADOOP_HOME}/share/hadoop/tools/lib/hadoop-streaming-2.2.0.jar\
    -D mapred.reduce.tasks=2\
    -D mapreduce.iterator.no=100\
    -files mapper, reducer\
    -input $INPUT_PATH\
    -output $OUTPUT_PATH\
    -mapper mapper\
    -reducer reducer
```

# Streaming程序运行方式



```
root@SY-0266:/opt/pgs/yarn-client# bin/hadoop jar share/hadoop/tools/lib/hadoop-streaming-2.2.0.jar -info
Usage: $HADOOP_PREFIX/bin/hadoop jar hadoop-streaming.jar [options]
Options:
  -input          <path> DFS input file(s) for the Map step.
  -output         <path> DFS output directory for the Reduce step.
  -mapper         <cmd|JavaClassName> Optional. Command to be run as mapper.
  -combiner       <cmd|JavaClassName> Optional. Command to be run as combiner.
  -reducer        <cmd|JavaClassName> Optional. Command to be run as reducer.
  -file          <file> Optional. File/dir to be shipped in the Job jar file.
                 Deprecated. Use generic option "-files" instead.
  -inputformat    <TextInputFormat(default)|SequenceFileAsTextInputFormat|JavaClassName>
                 Optional. The input format class.
  -outputformat   <TextOutputFormat(default)|JavaClassName>
                 Optional. The output format class.
  -partitioner    <JavaClassName> Optional. The partitioner class.
  -numReduceTasks <num> Optional. Number of reduce tasks.
  -inputreader    <spec> Optional. Input recordreader spec.
  -cmdenv         <n>=<v> Optional. Pass env.var to streaming commands.
  -mapdebug       <cmd> Optional. To run this script when a map task fails.
  -reduceddebug   <cmd> Optional. To run this script when a reduce task fails.
  -io            <identifier> Optional. Format to use for input to and output
                 from mapper/reducer commands
  -lazyOutput     Optional. Lazily create Output.
  -background     Optional. Submit the job and don't wait till it completes.
  -verbose       Optional. Print verbose output.
  -info          Optional. Print detailed usage.
  -help          Optional. Print help message.
```

The general command line syntax is  
bin/hadoop command [genericOptions] [commandOptions]



# Streaming程序运行方式说明



- 区分通用参数和命令行参数，通用参数应放在命令行参数前面，否则不起作用。通用参数有7个：
  - ✓ `-conf -D -fs -jt -files - libjars -archives`
- “`-file`” 或者 “`-files`” 参数，设置要分发到各个节点上的文件，对于mapper和reducer文件，必须要用或者 “`-files`” 或 “`-file`” 指定。
  - ✓ `-files mapper,reducer`
  - ✓ `-file mapper -file reducer`
- 每次运行程序前，清空输出目录  
`bin/hadoop fs -rmr /test/output`

## 实例2：用PHP实现Wordcount（Mapper实现）



```
#!/usr/bin/php
<?php
error_reporting(E_ALL ^ E_NOTICE);
$word2count = array();
// 标准输入为STDIN (standard input)
while (($line = fgets(STDIN)) !== false) {
    // 移除空白
    $line = trim($line);
    // 将行拆解成若干个单词
    $words = preg_split('/\W/', $line, 0, PREG_SPLIT_NO_EMPTY);
    // 将结果写到 STDOUT (standard output)
    foreach ($words as $word) {
        // 印出 [字 , "tab字符" , "数字" , "结束符"]
        echo $word, chr(9), "1", PHP_EOL;
    }
}
?>
```

## 实例2：用PHP实现Wordcount（Reducer实现）



```
#!/usr/bin/php
<?php
error_reporting(E_ALL ^ E_NOTICE);
$word2count = array();
// 标准输入为 STDIN
while (($line = fgets(STDIN)) !== false) {
    // 移除多余空白
    $line = trim($line);
    // 每一行的格式为(单词 "tab" 数字), 存入($word, $count)
    list($word, $count) = explode(chr(9), $line);
    // 转换格式string -> int
    $count = intval($count);
    //汇总
    $word2count[$word] += $count;
}
// 将结果写到 STDOUT (standard output)
foreach ($word2count as $word => $count) {
    echo $word, chr(9), $count, PHP_EOL;
}
?>
```

# 实例2：用PHP实现Wordcount（测试运行）



## ●测试mapper和reducer

✓ cat test.txt| php mapper.php | sort | php reducer.php

## ●在Hadoop上运行

```
#!/bin/bash
HADOOP_HOME=/opt/pgs/yarn-client
INPUT_PATH=/test/input
OUTPUT_PATH=/test/output
echo "Clearing output path: $OUTPUT_PATH"
$HADOOP_HOME/bin/hadoop fs -rmr $OUTPUT_PATH

${HADOOP_HOME}/bin/hadoop jar \
    ${HADOOP_HOME}/share/hadoop/tools/lib/hadoop-streaming-2.2.0.jar \
    -files mapper.php,reducer.php \
    -input $INPUT_PATH \
    -output $OUTPUT_PATH \
    -mapper "php mapper.php" \
    -reducer "php reducer.php" \
```

# 实例3：用Shell脚本语言实现Wordcount (Mapper实现)



```
#!/bin/bash
while read LINE; do
    for word in $LINE
    do
        echo "$word 1"
    done
done
```

# 实例3：用Shell脚本语言实现Wordcount (Reducer实现)



```
#!/bin/bash
count=0
started=0
word=""
while read LINE;do
    newword=`echo $LINE | cut -d ' ' -f 1`
    if [ "$word" != "$newword" ];then
        [ $started -ne 0 ] && echo "$word\t$count"
        word=$newword
        count=1
        started=1
    else
        count=$(( $count + 1 ))
    fi
done
echo "$word\t$count"
```

# 实例3：用Shell脚本语言实现Wordcount（测试运行）



## ●测试mapper和reducer

✓ `cat test.txt | sh mapper.sh | sort | sh reducer.sh`

## ●在Hadoop上运行

```
#!/bin/bash
HADOOP_HOME=/opt/pgs/yarn-client
INPUT_PATH=/test/input
OUTPUT_PATH=/test/output
echo "Clearing output path: $OUTPUT_PATH"
$HADOOP_HOME/bin/hadoop fs -rmr $OUTPUT_PATH

${HADOOP_HOME}/bin/hadoop jar\
    ${HADOOP_HOME}/share/hadoop/tools/lib/hadoop-streaming-2.2.0.jar\
    -files mapper.sh,reducer.sh\
    -input $INPUT_PATH\
    -output $OUTPUT_PATH\
    -mapper "sh mapper.sh"\
    -reducer "sh reducer.sh"
```



## ● 定义Hadoop Counter

- ✓ 使用标准错误输出，格式为

reporter:counter:<group>,<counter>,<amount>

## ● 在运行过程中展示状态信息

- ✓ 使用标准错误输出，格式为

reporter:status:<message>

## ● 打印调试信息

- ✓ 使用标准错误输出，调试信息将被保存stderr文件中

```
string key;
while(cin >> key) {
    cout << key << "\t" << "1" << endl;
    // Define counter named counter_no in group counter_group
    cerr << "reporter:counter:counter_group,counter_no,1\n";
    // display status
    cerr << "reporter:status:processing.....\n";
    // Print logs for testing
    cerr << "This is log, will be printed in stdout file\n";
}
return 0;
```

## ● 获取conf配置信息，比如reduce task个数等

- ✓ 在环境变量中获取，所有的“.”被替换成了“\_”

```
int main(int argc, char *argv[], char *env[]) {
```

比如：“mapreduce.job.reduces”被换成了“mapreduce\_job\_reduces”





## ● 对多字段文本数据的支持

```
$HADOOP_HOME/bin/hadoop jar $HADOOP_HOME/hadoop-streaming.jar \  
-D stream.map.output.field.separator=. \  
-D stream.num.map.output.key.fields=4 \  
-D map.output.key.field.separator=. \  
-D mapred.text.key.partitionner.options=-k1,2 \  
-D mapred.reduce.tasks=12 \  
-input myInputDirs \  
-output myOutputDir \  
-mapper org.apache.hadoop.mapred.lib.IdentityMapper \  
-reducer org.apache.hadoop.mapred.lib.IdentityReducer \  
-partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner
```

```
11.12.1.2  
11.14.2.3  
11.11.4.1  
11.12.1.1  
11.14.2.2
```

```
11.11.4.1  
-----  
11.12.1.2  
11.12.1.1  
-----  
11.14.2.3  
11.14.2.2
```

```
11.11.4.1  
-----  
11.12.1.1  
11.12.1.2  
-----  
11.14.2.2  
11.14.2.3
```

```
$HADOOP_HOME/bin/hadoop jar $HADOOP_HOME/hadoop-streaming.jar \  
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedComparator \  
-D stream.map.output.field.separator=. \  
-D stream.num.map.output.key.fields=4 \  
-D map.output.key.field.separator=. \  
-D mapred.text.key.comparator.options=-k2,2nr \  
-D mapred.reduce.tasks=12 \  
-input myInputDirs \  
-output myOutputDir \  
-mapper org.apache.hadoop.mapred.lib.IdentityMapper \  
-reducer org.apache.hadoop.mapred.lib.IdentityReducer
```

```
11.12.1.2  
11.14.2.3  
11.11.4.1  
11.12.1.1  
11.14.2.2
```

```
11.14.2.3  
11.14.2.2  
11.12.1.2  
11.12.1.1  
11.11.4.1
```

# Java与Streaming编程方式比较



- **Java编程**

- ✓ **Hadoop最原始开发语言；**
- ✓ **支持所有功能，是其他编程方式的基础。**

- **Streaming编程**

- ✓ **仅用于开发Mapper和Reducer，其他组件需采用Java实现；**
- ✓ **天生支持文本格式，但二进制格式支持较弱；**
- ✓ **通常用于简单的文本数据处理，加快开发效率。**



- **MapReduce的编程模型**
- **MapReduce编程接口介绍**
- **Java编程**
- **多语言编程**